

Part 3 and Part 4 Report: Query Implementation and performance optimization

- 1) Fetching all students enrolled in a specific course.

student_id	first_name	last_name
104	Deanna	Wolf
174	Abigail	Walker
277	Alejandro	Gill
386	Caitlin	Hull
483	Melvin	Wallace
104	Deanna	Wolf
174	Abigail	Walker
277	Alejandro	Gill
386	Caitlin	Hull
483	Melvin	Wallace

- 2) Calculating the average number of students enrolled in courses offered by a particular instructor at the university.

avg_students
3.3333333333333335

3) Listing all courses offered by a specific department.

course_id	course_name
1	Leader thus.
3	Someone.
13	Value fine subject.
16	Close.
27	Run.
30	Person behind.
36	Sign kind.
41	Home.
44	Clearly office.
53	Under share.
71	War short since a...
84	Modern speak smal...
124	Seat tonight.
126	Mr treat ever enter.
128	Science bill line.
139	Beautiful design ...
1	Leader thus.
3	Someone.
13	Value fine subject.
16	Close.

only showing top 20 rows

4) Finding the total number of students per department.

department_name	total_students
Brown LLC	232
Schwartz Inc	172
Aguilar, Taylor a...	228
Campbell Ltd	152
Santos Ltd	192
Munoz-Johnson	240
Carter LLC	168
Young-O'Brien	180
Clayton-Johnson	212
Green Group	224

- 5) Finding instructors who have taught all the BTech CSE core courses sometime during their tenure at the university.

instructor_id	first_name	last_name
1	Pamela	Ashley
5	Noah	Patterson
2	Edward	Hurst

- 6) Finding top-10 courses with the highest enrollments.

course_id	course_name	enrollments
93	Less dream image ...	7
93	Less dream image ...	7
104	State although.	7
135	Discussion interv...	7
136	Become even.	7
104	State although.	7
135	Discussion interv...	7
136	Become even.	7
117	Small guess.	6
128	Science bill line.	6

1) Indexing Optimisation

- `db.students.createIndex({ "student_id": 1 })`
 - `db.courses.createIndex({ "course_id": 1 })`
 - `db.departments.createIndex({ "department_id": 1 })`
 - `db.instructors.createIndex({ "instructor_id": 1 })`
 - `db.courses.createIndex({ "course_id": 1, "instructor_id": 1 })`
-
- Indexes allow MongoDB to quickly locate and access relevant documents without scanning the entire collection. By indexing fields like **student_id**, **course_id**, and **instructor_id**, the database can reduce search times and lower input/output operations, leading to faster data retrieval.
 - Compound indexes (e.g., on `course_id` and `instructor_id`) improve performance for queries involving multiple fields, reducing the need for post-retrieval filtering. This is crucial for maintaining query efficiency as the database grows, ensuring consistent performance in read-heavy applications.

Query Number	Execution_Time_Without_Optimization	Execution_Time_Indexing
1	0.0745 seconds	0.0672 seconds
2	0.0380 seconds	0.1107 seconds
3	0.0561 seconds	0.0611 seconds
4	0.0493 seconds	0.1130 seconds
5	0.1789 seconds	0.1513 seconds
6	0.0915 seconds	0.0587 seconds

2) Partitioning Optimization

By adjusting the number of partitions, we can ensure that our workload is evenly distributed across the various runners. This helps prevent some runners from getting over burdened than others.

i) On basis of dataframe

- `students_df = students_df.repartition(n_partitions)`
- `departments_df = departments_df.repartition(n_partitions)`
- `courses_df = courses_df.repartition(n_partitions)`
- `instructors_df = instructors_df.repartition(n_partitions)`

ii) On basis of Column

- `students_df = students_df.repartition(n_partitions, "student_id")`
- `departments_df = departments_df.repartition(n_partitions, "department_id")`
- `courses_df = courses_df.repartition(n_partitions, "course_id")`
- `instructors_df = instructors_df.repartition(n_partitions, "instructor_id")`

Query Number	Execution_Time_Without_Optimization	Execution_Time_Partitioning_df	Execution_Time_Partitioning_column
1	0.0745 seconds	0.0366 seconds	0.1371 seconds
2	0.0380 seconds	0.0853 seconds	0.1107 seconds
3	0.0561 seconds	0.0219 seconds	0.0526 seconds
4	0.0493 seconds	0.1459 seconds	0.0907 seconds
5	0.1789 seconds	0.1905 seconds	0.1044 seconds
6	0.0915 seconds	0.0398 seconds	0.0772 seconds

3) Partitioning Optimization and Indexing.

Query Number	Execution_Time_Without_Optimization	Execution_Time_Partitioning_and_Indexing
1	0.0745 seconds	0.3370 seconds
2	0.0380 seconds	0.2067 seconds

3	0.0561 seconds	0.0714 seconds
4	0.0493 seconds	0.1489 seconds
5	0.1789 seconds	0.0533 seconds
6	0.0915 seconds	0.0587 seconds

- These optimizations will show a much more significant reduction in execution time when applied to larger datasets or in a distributed computing environment.
- For smaller datasets, the overhead of creating indexes or repartitioning can sometimes result in increased query execution times.